

# Orchestrating Azure Container Service Cluster using Kubernetes

---

SONU SATHYADAS



sonusathyadas@hotmail.com



<https://in.linkedin.com/in/sonusathyadas>



@sonusathyadas

# Sonu Sathyadas

---

Tech Lead & Training Consultant @ Synergetics India

8+ Years of corporate training experience



## Microsoft

Specialist

---

Developing Microsoft  
Azure Solutions

## Microsoft

Specialist

---

Implementing Microsoft  
Azure Infrastructure  
Solutions

# Agenda

---

- Containerization
- Docker
- Docker Hub and Azure Container Registry
- Kubernetes
- Azure Container Services
- Conclusion

# Containers

---

OS-level virtualization method for deploying and running distributed applications.

Run on a single control host and access a single kernel.

# Hardware Virtualization Vs Application Isolation

---

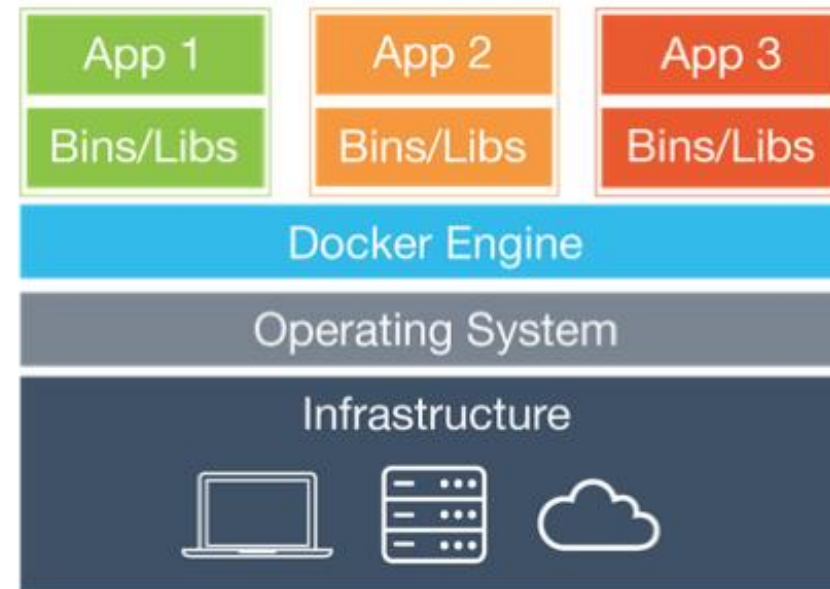
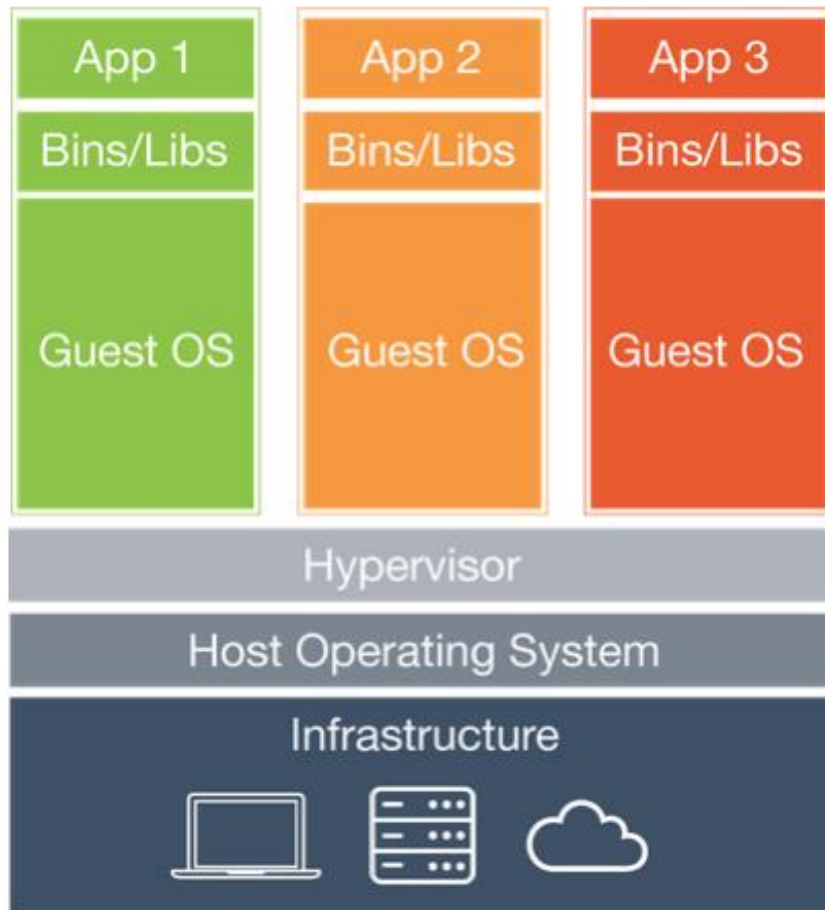
- Containers and virtual machines both allow you to abstract the workload from the underlying hardware.
- VMs mimics a complete server - OS, Applications, Network.
- VMs runs on top of Hypervisor (eg: VMWare, Hyper-V)
- Containerization allows virtual instances to share a single host operating system and relevant binaries, libraries or drivers.
- Containerization is handled by a containerization engine, like Docker.

# Virtual Machines Vs Containers

---

- Containers are an abstraction at the app layer that packages code and dependencies together.
- Multiple containers can run on the same machine and share the OS kernel with other containers, each running as isolated processes in user space.
- Containers take up less space than VMs
- Containers size is small (MBs)
- Provisioning containers only take a few seconds or less.
- Containers are a very cost effective solution.
- Container-based virtualization are a great option for microservices, DevOps, and continuous deployment.
- Virtual machines are an abstraction of physical hardware turning one server into many servers
- VM hypervisor allows multiple VMs to run on a single machine.
- Each VM includes a full copy of an operating system, one or more apps, necessary binaries and libraries - taking up tens of GBs.
- VMs can also be slow to boot.

# Virtual Machines Vs Containers



# Advantages of containers

---

Container is notably smaller, easier to migrate or download, faster to backup or restore and requires less memory.

No licence cost.

Optimized utilization of host resources.

Support many more containers on the same infrastructure.

Fast deployment than VMs

Better for microservices and continuous integration and delivery.



# Disadvantages of Containers

---

Lack of isolation from the host OS

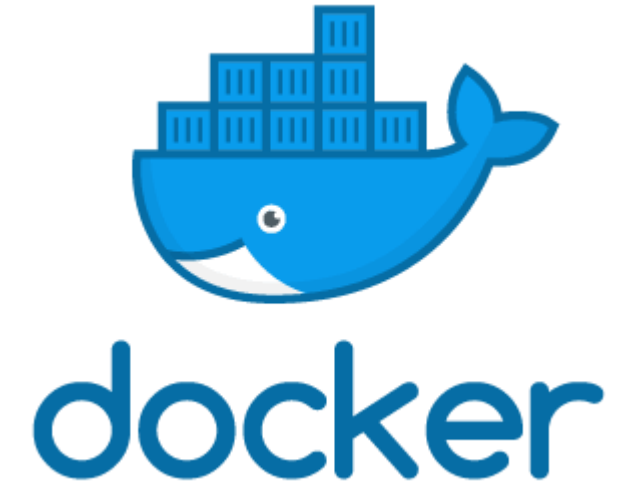
Each container must use the same OS as the base OS.

Lack of security - Host OS attack affect containers too.

# Container implementations

---

- Docker
- Rocket (rkt) from CoreOS
- LXD containerization engine for Ubuntu
- BSD Jails
- LXC
- Solaris Zones
- NonLinux-based Container Solutions
  - Windows Server containers and Hyper-V containers



# Docker

---

IMPLEMENTING DOCKER CONTAINERS

# Docker images

---

Docker is the company driving the container movement and the only container platform provider to address every application across the hybrid cloud.

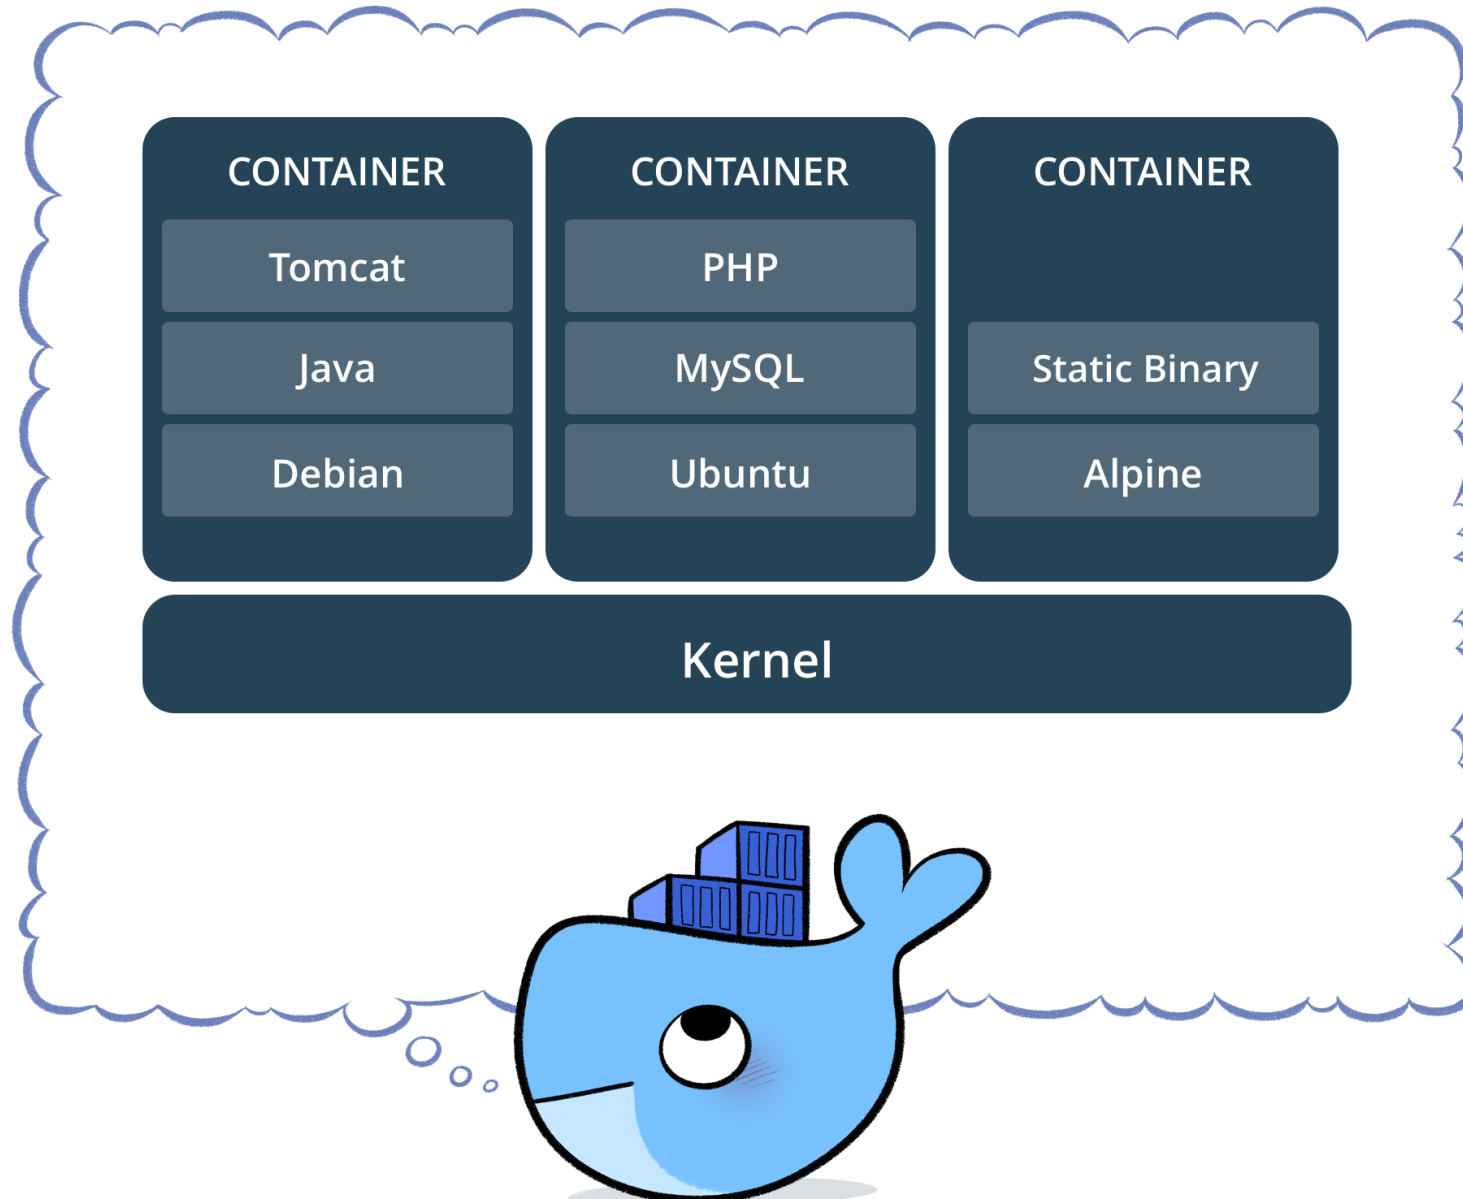
Container image is a lightweight, stand-alone, executable package of a piece of software.

Docker image includes everything needed to run it: code, runtime, system tools, system libraries, settings.

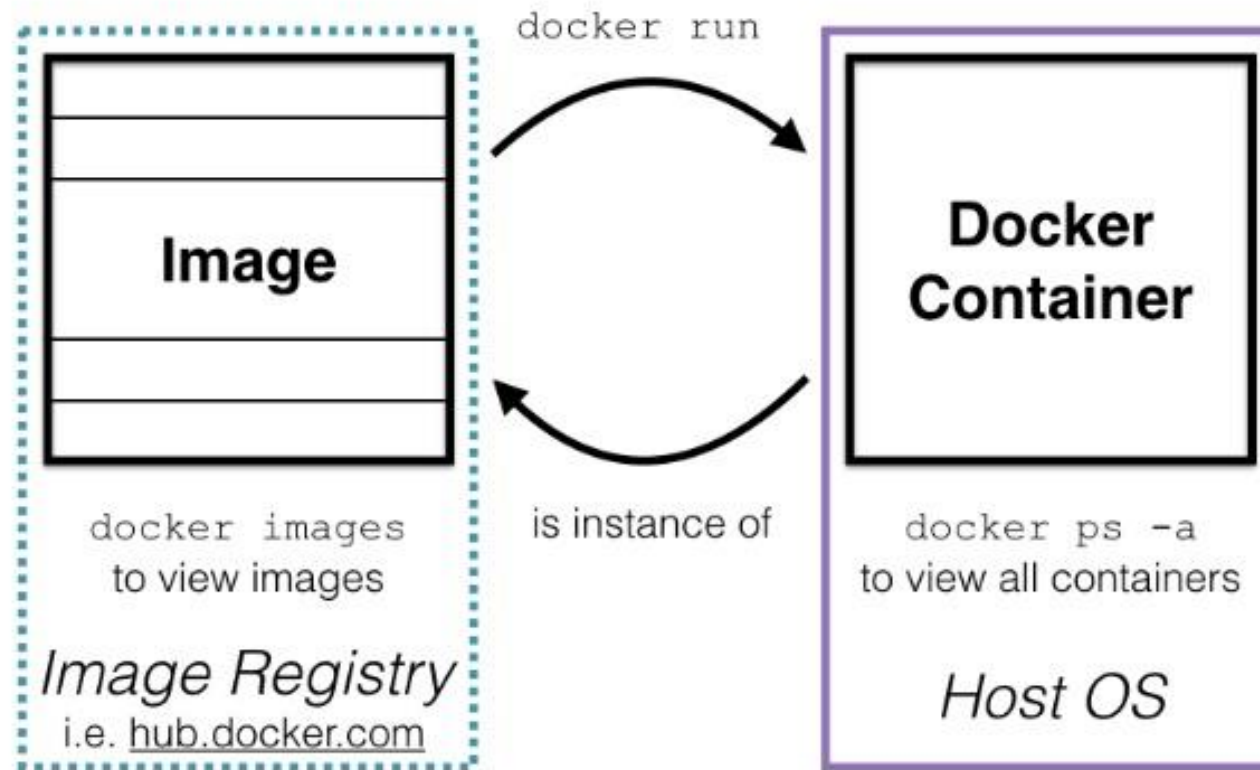
Available for both Linux and Windows based apps

Containers isolate software from its surroundings

Images are constructed from filesystem layers and share common files to minimize disk usage and image downloads are much faster.



# Docker images VS Containers

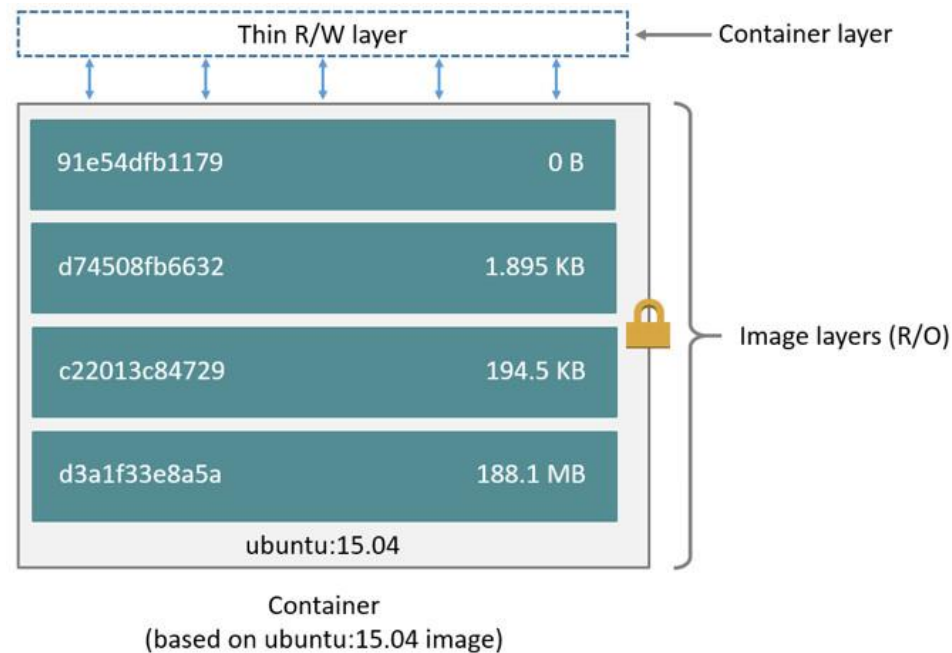


# Docker image and layers

Docker image is built up from a series of layers.

Each layer represents an instruction in the image's Dockerfile.

Each layer except the very last one is read-only.



# Dockerfile

---

A Dockerfile is a text document that contains all the commands a user could call on the command line to assemble an image.

```
FROM ubuntu:15.04  
COPY . /app  
RUN make /app  
CMD python /app/app.py
```

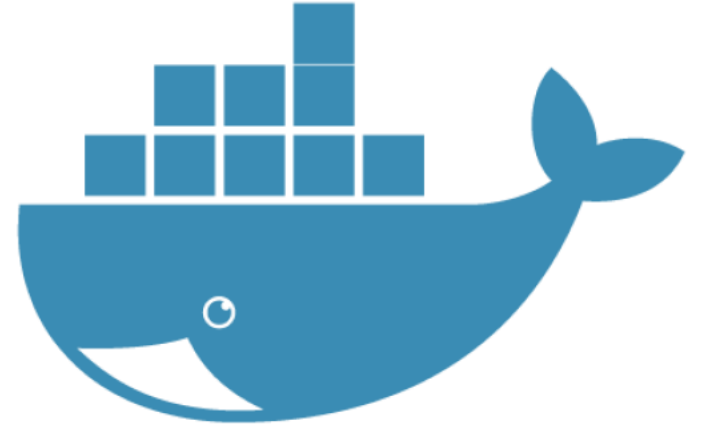


Command	Description
<code>docker build</code>	Build an image from a Dockerfile
<code>docker commit</code>	Create a new image from a container's changes
<code>docker cp</code>	Copy files/folders between a container and the local filesystem
<code>docker create</code>	Create a new container
<code>docker exec</code>	Run a command in a running container
<code>docker image</code>	Manage images
<code>docker images</code>	List images
<code>docker login</code>	Log in to a Docker registry
<code>docker logout</code>	Log out from a Docker registry
<code>docker ps</code>	List containers
<code>docker pull</code>	Pull an image or a repository from a registry
<code>docker push</code>	Push an image or a repository to a registry
<code>docker rename</code>	Rename a container
<code>docker restart</code>	Restart one or more containers
<code>docker rm</code>	Remove one or more containers
<code>docker rmi</code>	Remove one or more images
<code>docker run</code>	Run a command in a new container
<code>docker start</code>	Start one or more stopped containers
<code>docker stop</code>	Stop one or more running containers

# DEMO

---

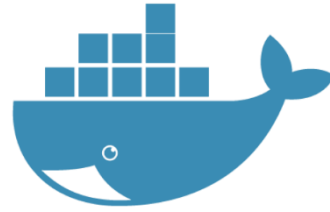
CREATING DOCKER CONTAINERS



# Docker Hub

---

THE REPOSITORY FOR DOCKER IMAGES



# Docker Hub

---

Repository for docker images.

Public or private repositories.

Organization in Docker Cloud contains Teams, and each Team contains users.

Push images to Docker hub

- Login to docker hub account
- Build the image with a tag
- Push images using '***Docker push***' command

Pull the image using '***docker pull***' command

[Dashboard](#)[Explore](#)[Organizations](#)[Create](#)[sonusathyadas](#)[sonusathyadas](#)[Repositories](#)[Stars](#)[Contributed](#)Private Repositories: Using 1 of 1 [Get more](#)

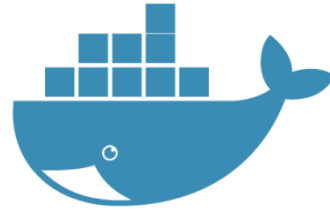
## Repositories

[Create Repository +](#)[sonusathyadas/products-api](#)  
public0  
STARS19  
PULLS>  
DETAILS[sonusathyadas/node-microservice](#)  
public0  
STARS3  
PULLS>  
DETAILS[sonusathyadas/office-management-app](#)  
private0  
STARS0  
PULLS>  
DETAILS

### Docker Security Scanning

Protect your repositories from vulnerabilities.

[Try it free](#)



# Push images to Docker Hub

---

Login to Docker Hub

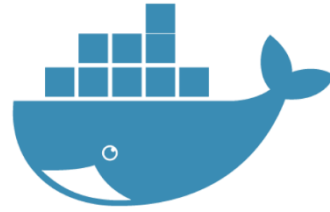
***docker login***

Build the image in local repository.

**docker build -t <imagename:tag> .**

Upload the docker image to Docker hub

**docker push <imagename:tag>**



# Pull image from Docker Hub

---

Login to Docker Hub

***docker login***

Pull image from Docker Hub

***docker pull <imagename:tag>***



# Plan and Pricing

Plan	Price	Private Repositories	Parallel Builds
Free	\$0/mo	1	1
Micro	\$7/mo	5	5
Small	\$12/mo	10	10
Medium	\$22/mo	20	20
Large	\$50/mo	50	50
XLarge	\$100/mo	100	100
XX-Large	\$250/mo	250	250
XXX-Large	\$500/mo	500	500





# Kubernetes

---

ORCHESTRATING CLUSTERS



# Kubernetes

---

- Initialized by Google
- Open-source orchestration tool for Clusters.
- Written in Go/GoLang
- Derived from Google's proprietary container management tools- Borg and omega
-



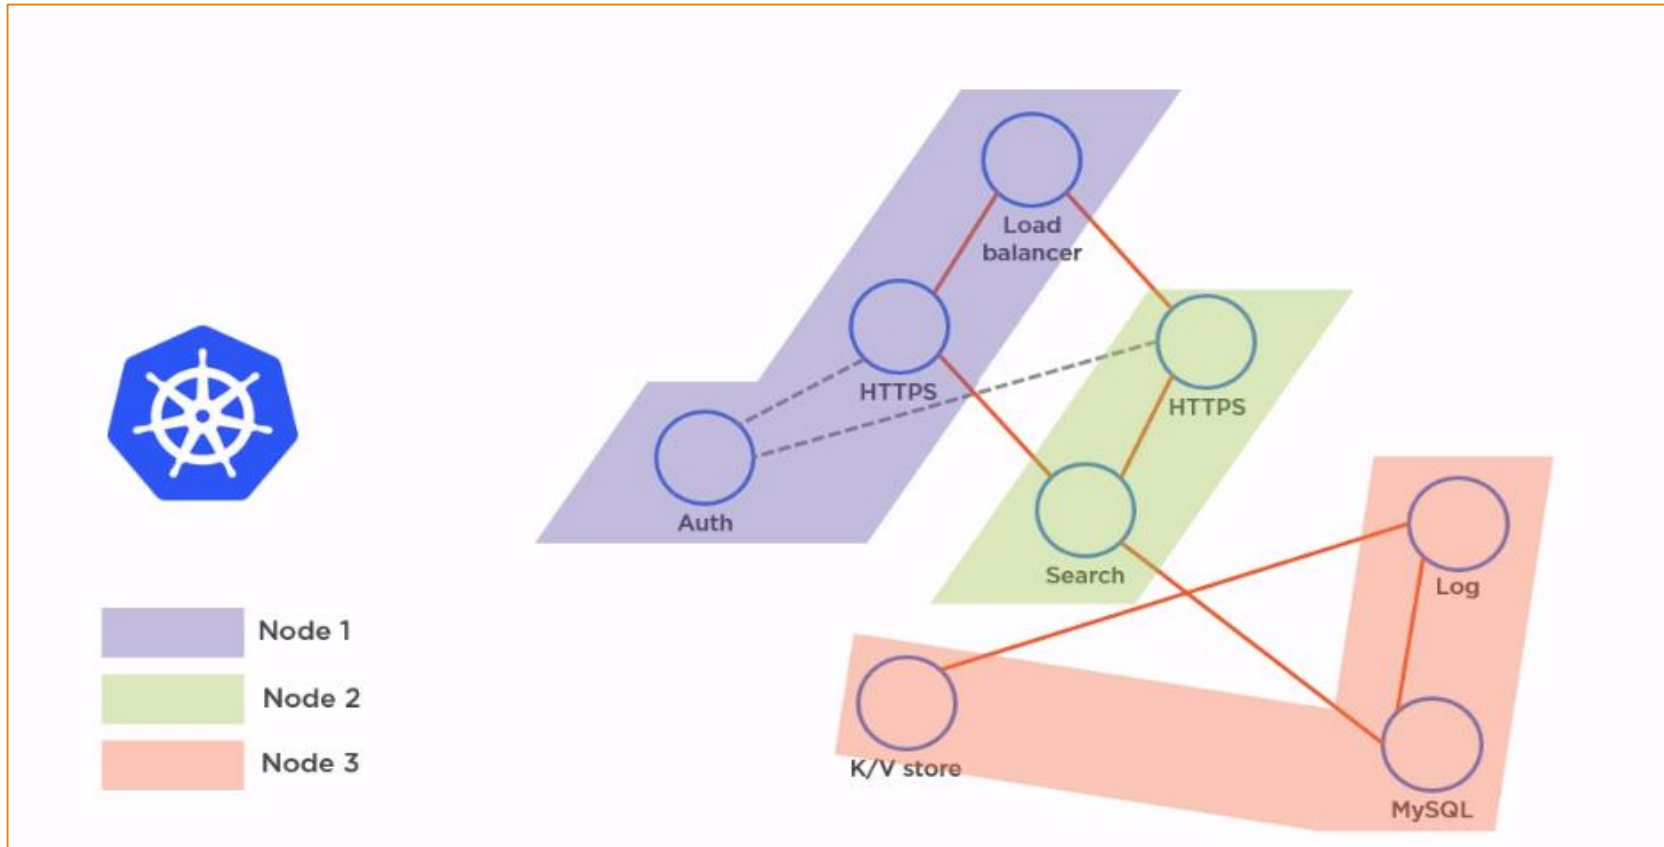
# Components

---

- Master
- Nodes
- Pod
- Replication Controller (ReplicaSet)
- Services
- Deployments



# Kubernetes - The Big picture



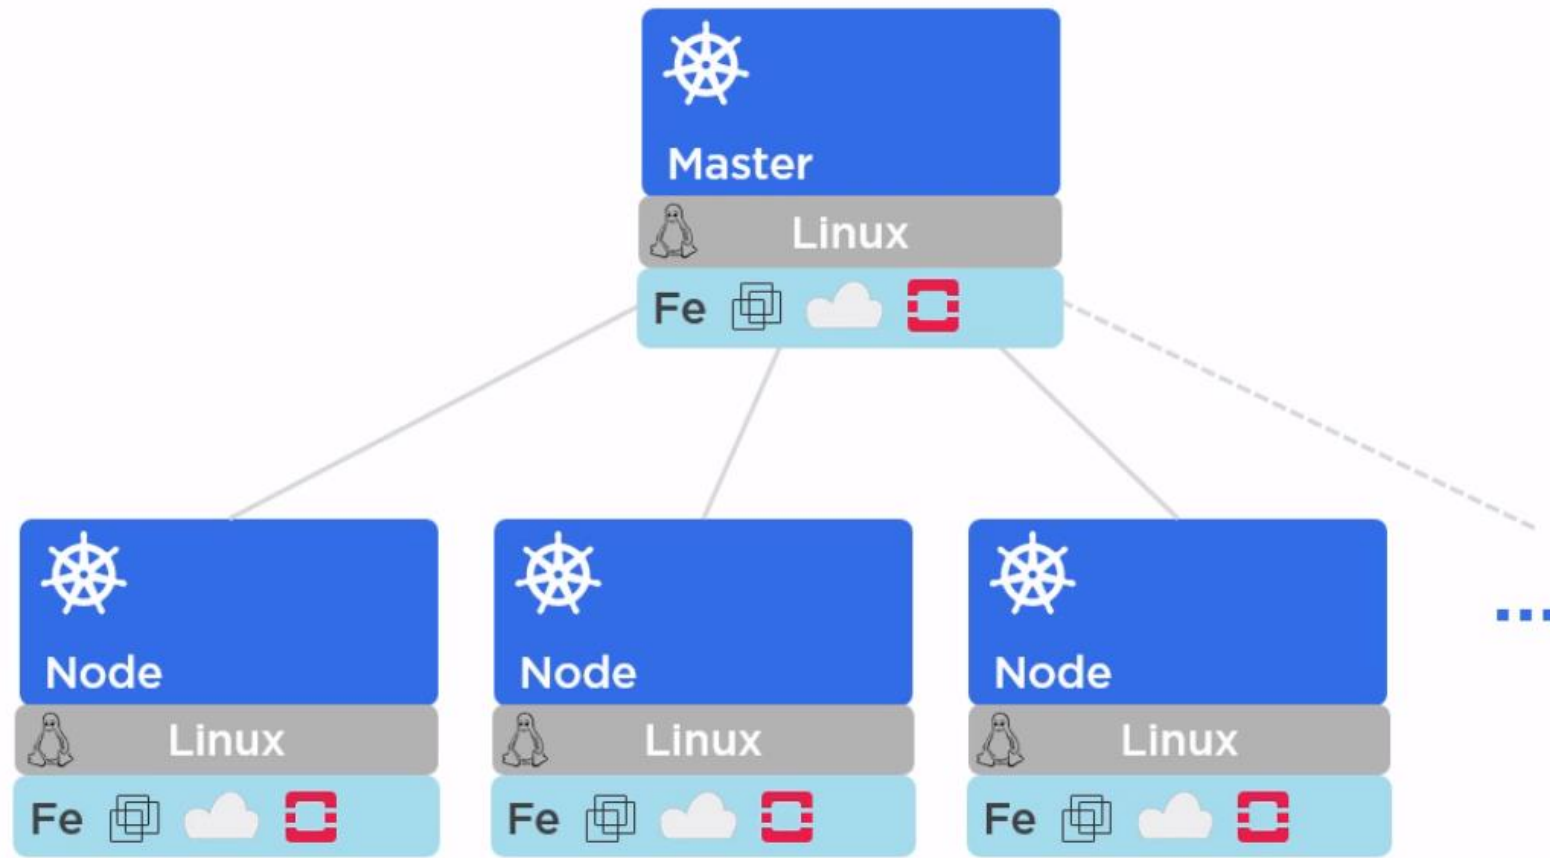
# Kubernetes Cluster

---

- Coordinates highly available cluster of machines that works as single unit.
- Machines – Master and Nodes
- Automates the distribution and scheduling of application containers across a cluster
- Apps need to be containerized.
- Containers runs on Pod
- Node contains multiple Pods
- Master controls Worker Nodes

# Kubernetes Cluster

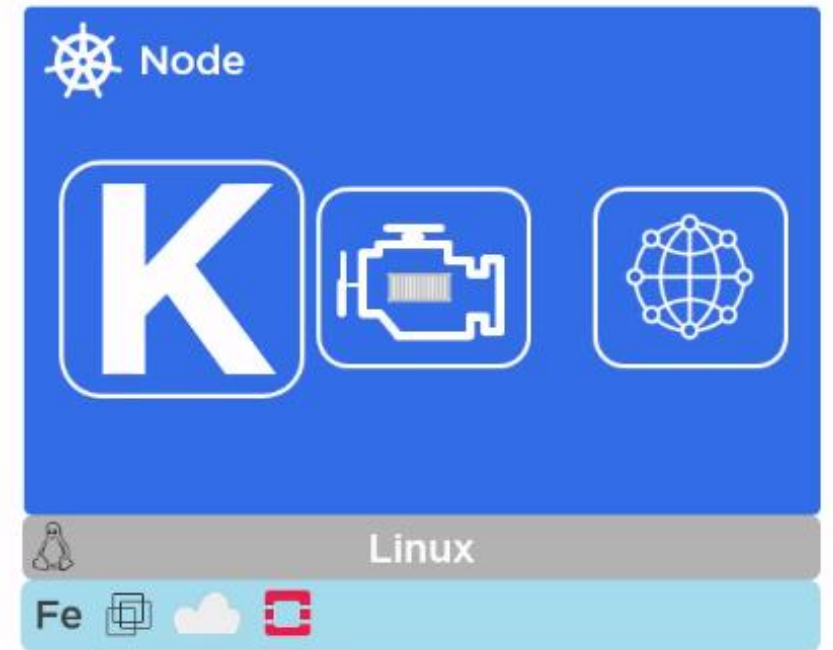
---



# Cluster Node

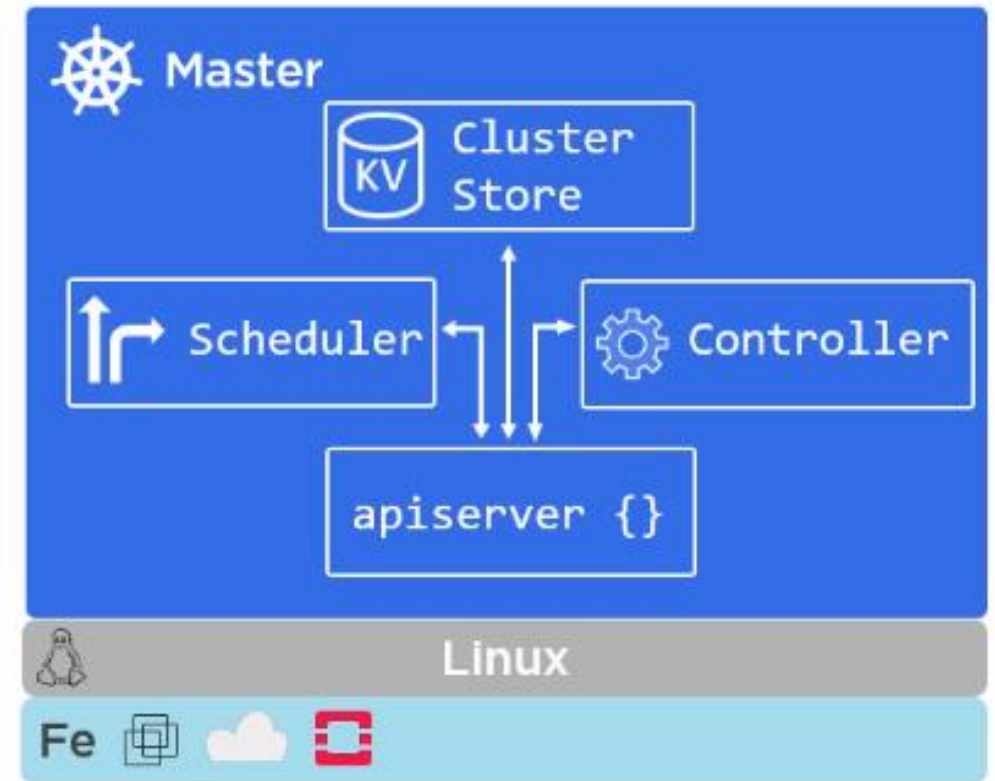
---

- Workers in a cluster
- Apps runs in node, not in master
- VM or Physical machine
- Node components
  - Kubelet
  - Container management tool - Docker or Rkt
- 3 node recommended for production (min).
- Communicates with master using kubernets API



# Cluster Master

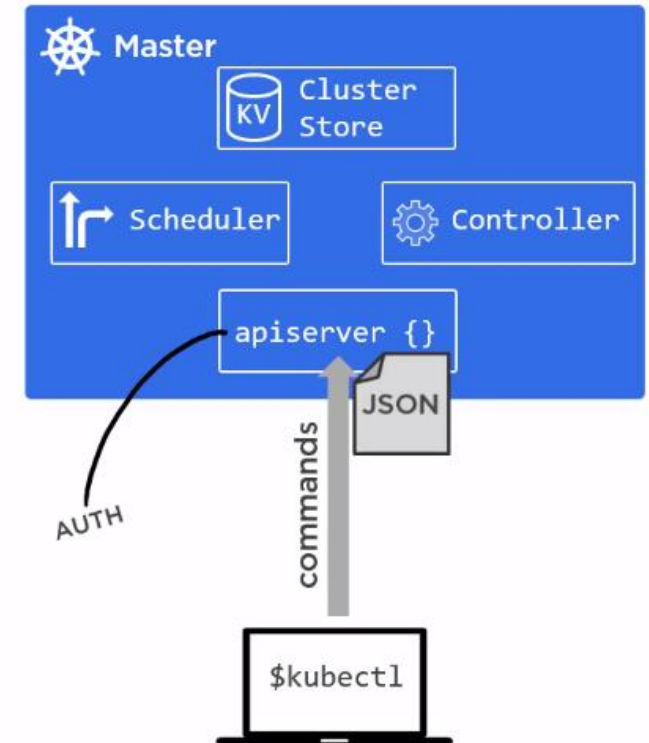
- Managing the cluster
- Coordinates all activities in cluster
  - Scheduling applications
  - Maintaining applications' desired state
  - Scaling applications
  - Rolling out new updates
- One or more masters for HA.
- Master components
  - Kube-apiserver
  - Cluster store (etcd)
  - Kube-controller-manager
  - Kube-scheduler





# Kube-apiserver

- Master are only exposed to public.
- Expose the REST API to connect with the cluster
- kube-apiserver exposes the Kubernetes API
- front-end for the Kubernetes control plane
- scale horizontally – scales by deploying more instances.



# Cluster Store (etcd)

---

- Used as Kubernetes' backing store.
- Always have a backup plan for etcd's data
- Stores Cluster state and configuration

# Kube-controller-manager

---

- Controller of controllers
- Runs background controller threads that handle routine tasks in the cluster
- Each controller is a separate process, but all are compiled into a single binary and run in a single process
- Controllers are
  - Node Controller
  - Replication Controller
  - Endpoints Controller
  - Service Account & Token Controllers

# Kube-scheduler

---

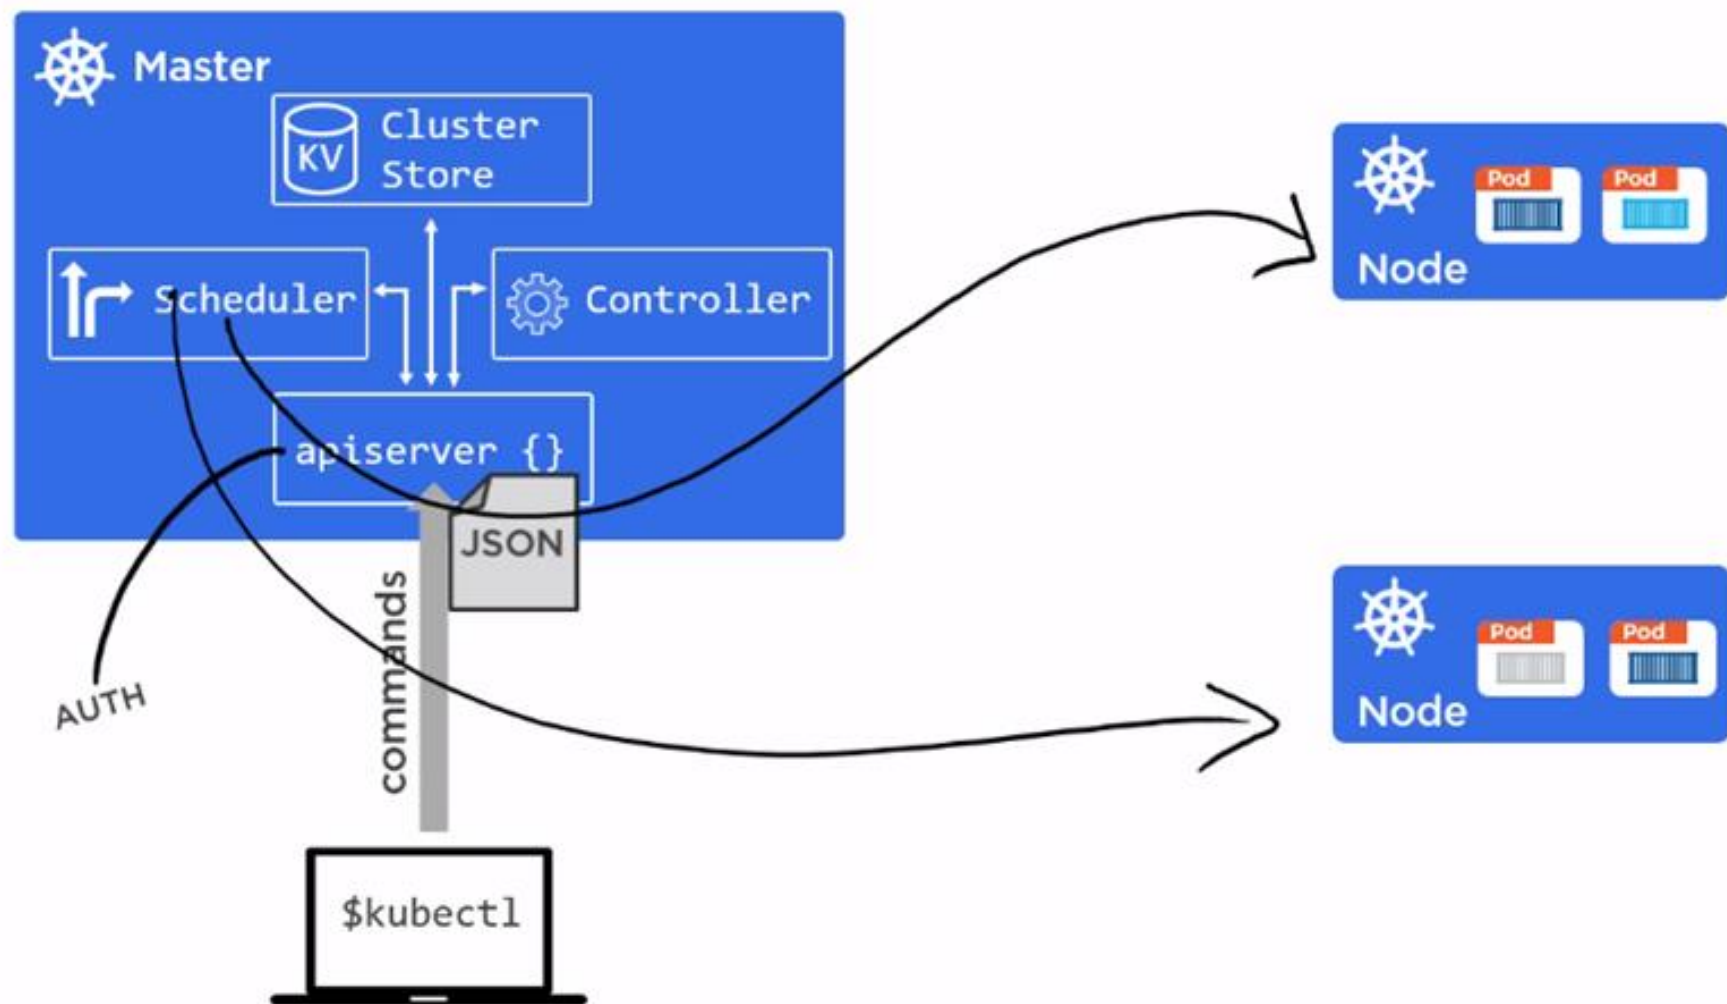
- Watches apiserver for new pods
- Assign works to nodes – affinity/anti-affinity, constraints, resources etc.
- Watches newly created pods that have no node assigned, and selects a node for them to run on


# Manifest file

- Declaration of deployments, Services, pods ...
- YAML or JSON

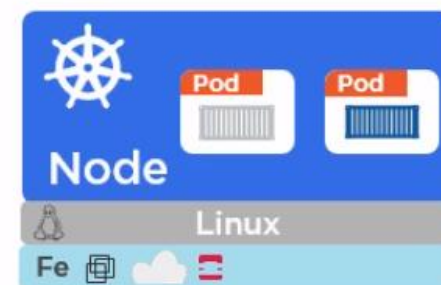
nginx-app.yaml 

```
apiVersion: v1
kind: Service
metadata:
  name: my-nginx-svc
  labels:
    app: nginx
spec:
  type: LoadBalancer
  ports:
    - port: 80
  selector:
    app: nginx
---
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: my-nginx
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```






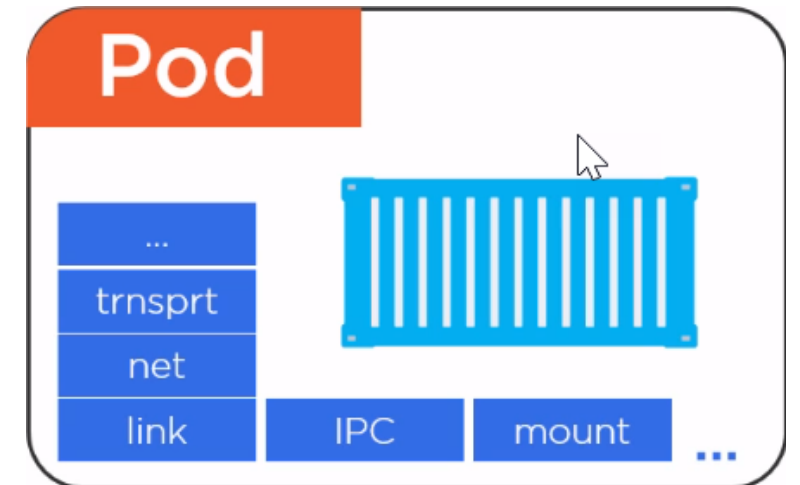
Don't run user workloads on  
"Master"



# Pod

---

- Basic building block of Kubernetes
- Smallest and simplest unit to create or deploy
- Pod represents a running process on your cluster
- Pod encapsulates
  - Application container (in some cases, multiple containers)
  - Storage resources
  - A unique network IP
  - Options that govern how the container(s) should run

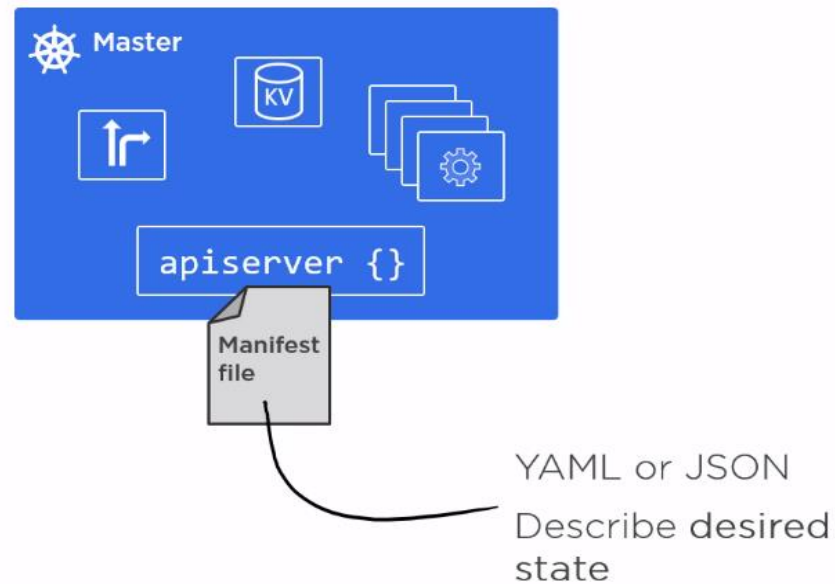


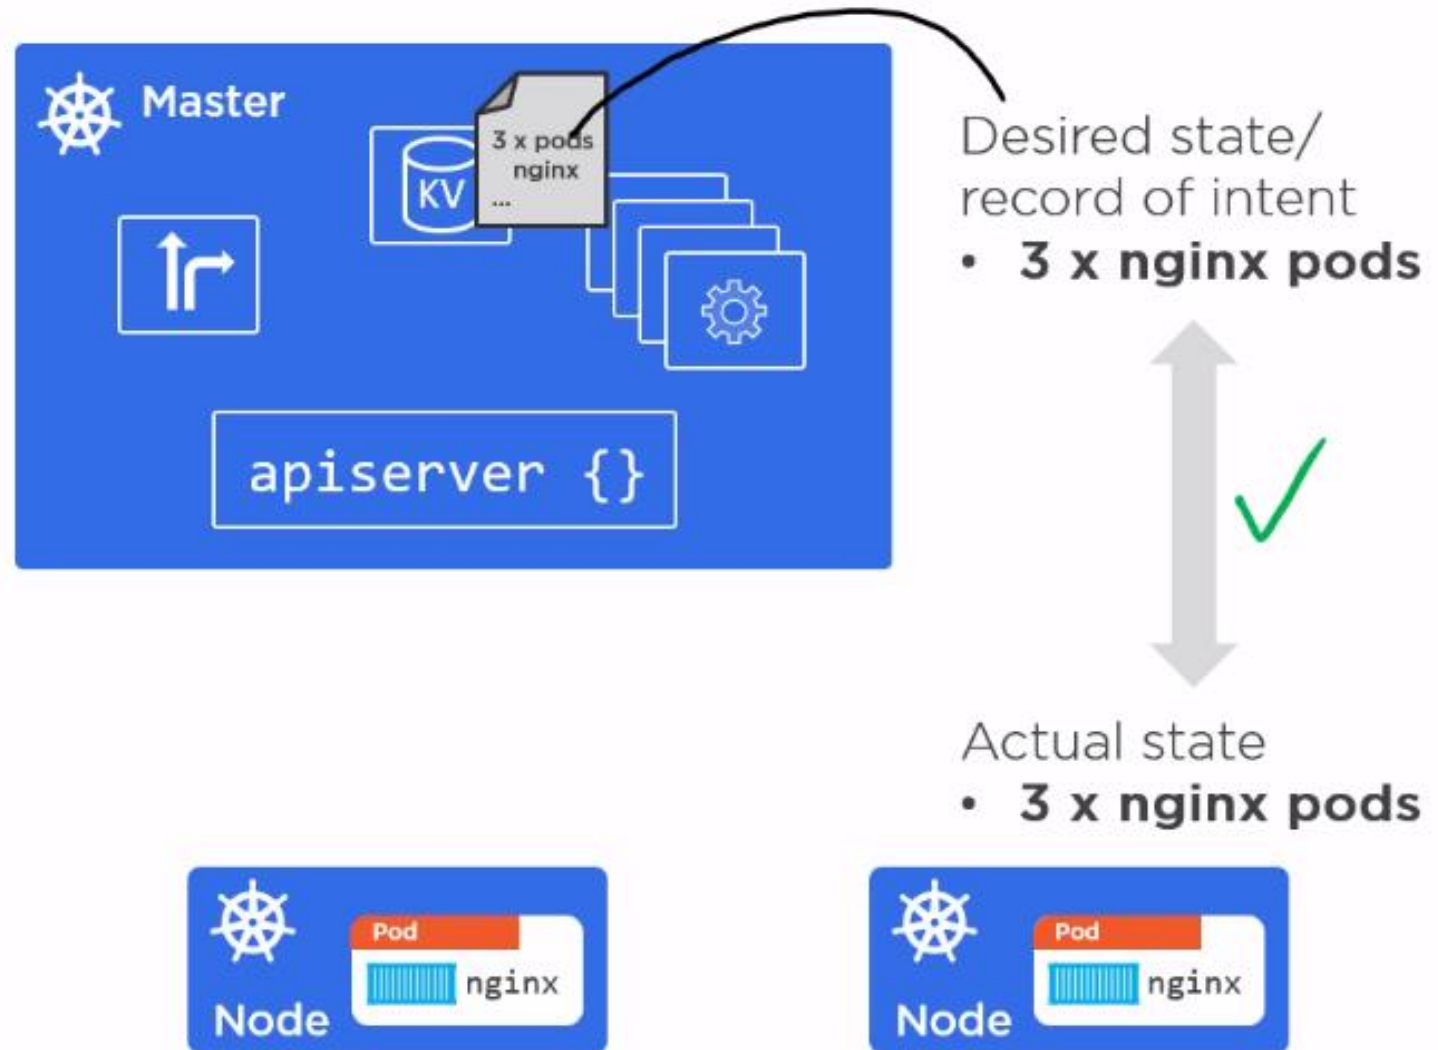


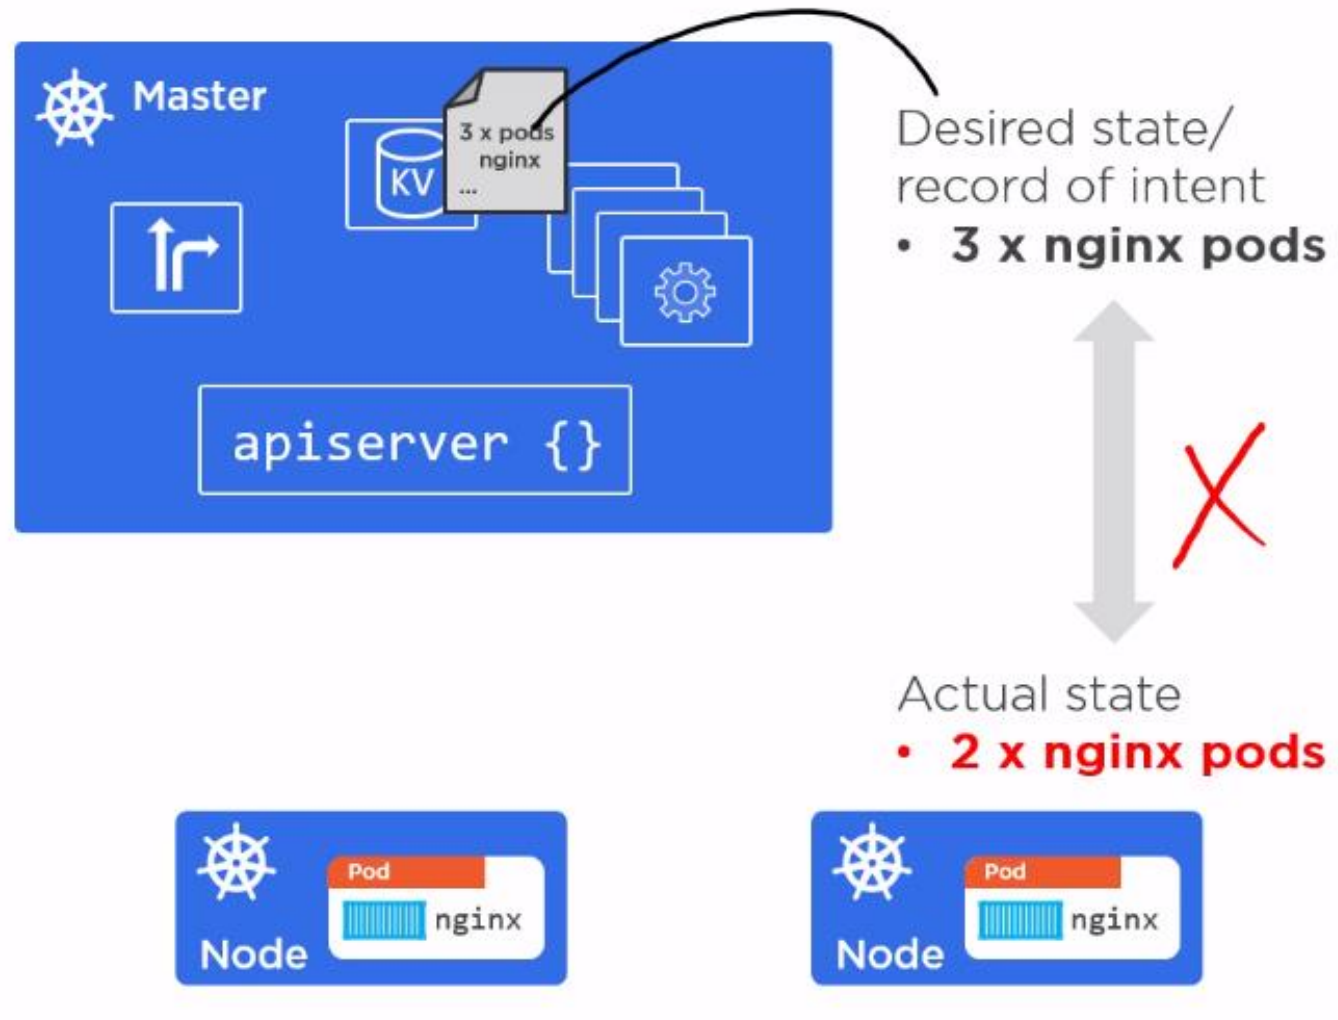
# Declarative Model and Desired State

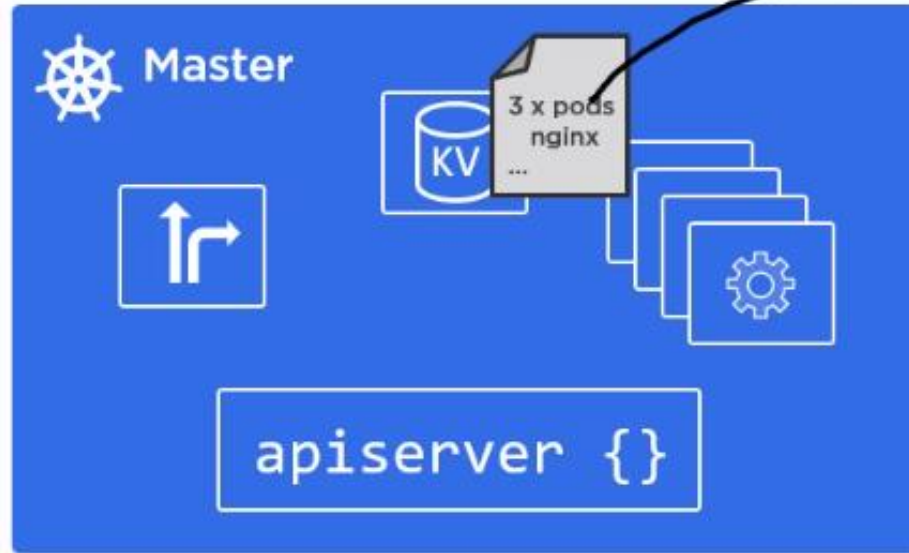
---

- Kubernetes objects can be created, updated, and deleted using configuration files.
- Configuration files are stored in a directory
- kubectl recursively create and update those objects as needed.









Desired state/  
record of intent

- **3 x nginx pods**



Actual state


- **3 x nginx pods**



# Pod scaling

---

Scale out is recommended



# Replication Controller (ReplicaSet)

---

Ensures a specified number of pod replicas are always up and available

ReplicaSet is the next-generation Replication Controller.

Only difference between a ReplicaSet and a Replication Controller is the selector support.

ReplicaSet supports the new set-based selector requirements

Replication Controller only supports equality-based selector requirements

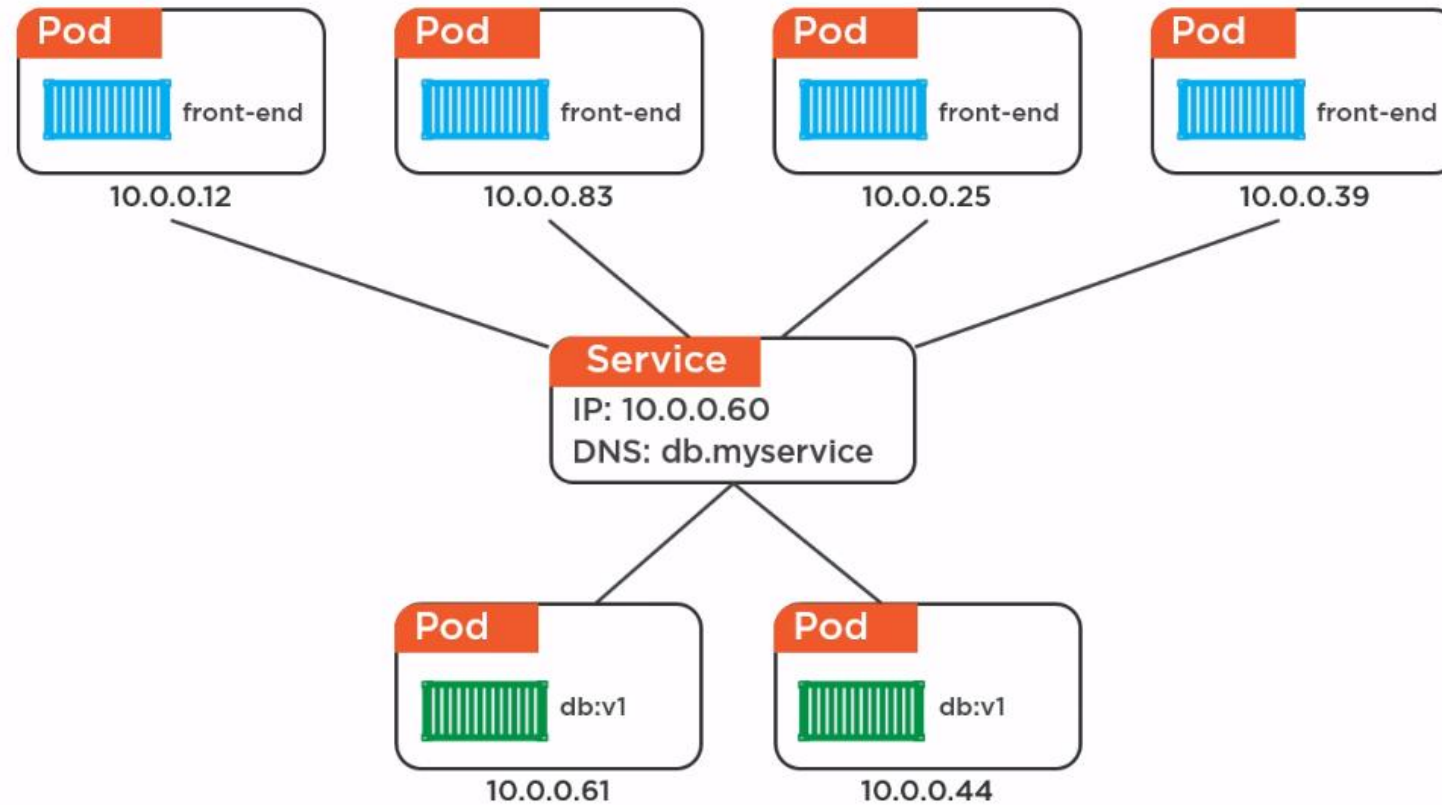
# Services

---

- An abstraction which defines a logical set of Pods and a policy by which to access them.
- Set of Pods targeted by a Service is usually determined by a Label Selector.
- Pods internal IP's are not reliable for communication.
- IP may change during scaling rolling deployment.
- Services acts as intermediary/Load balancer to access pods.

# Services

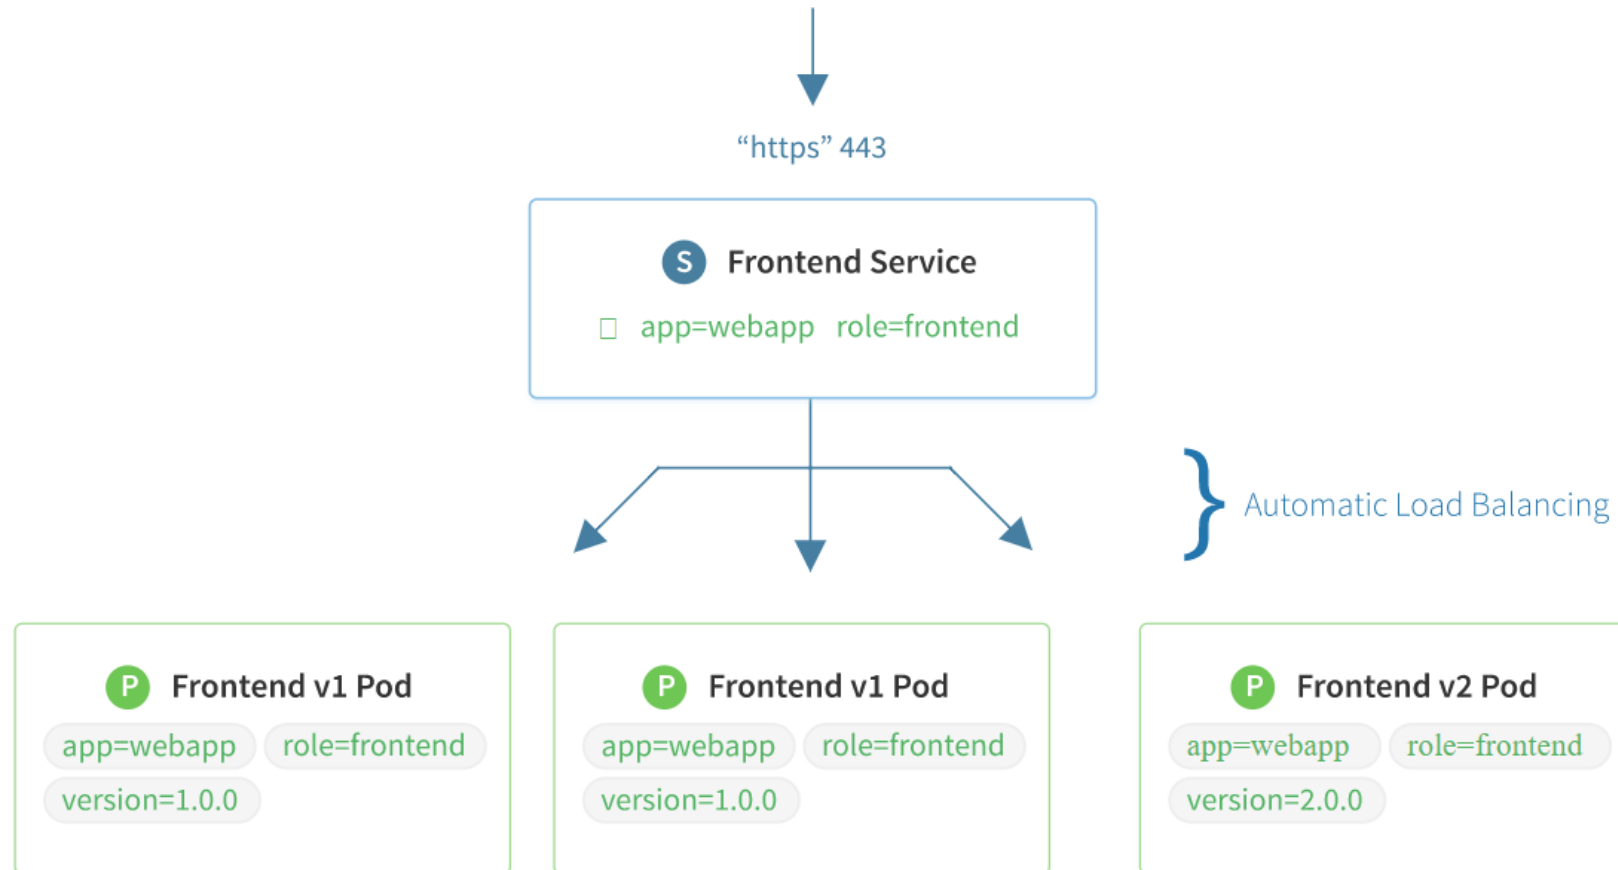
---



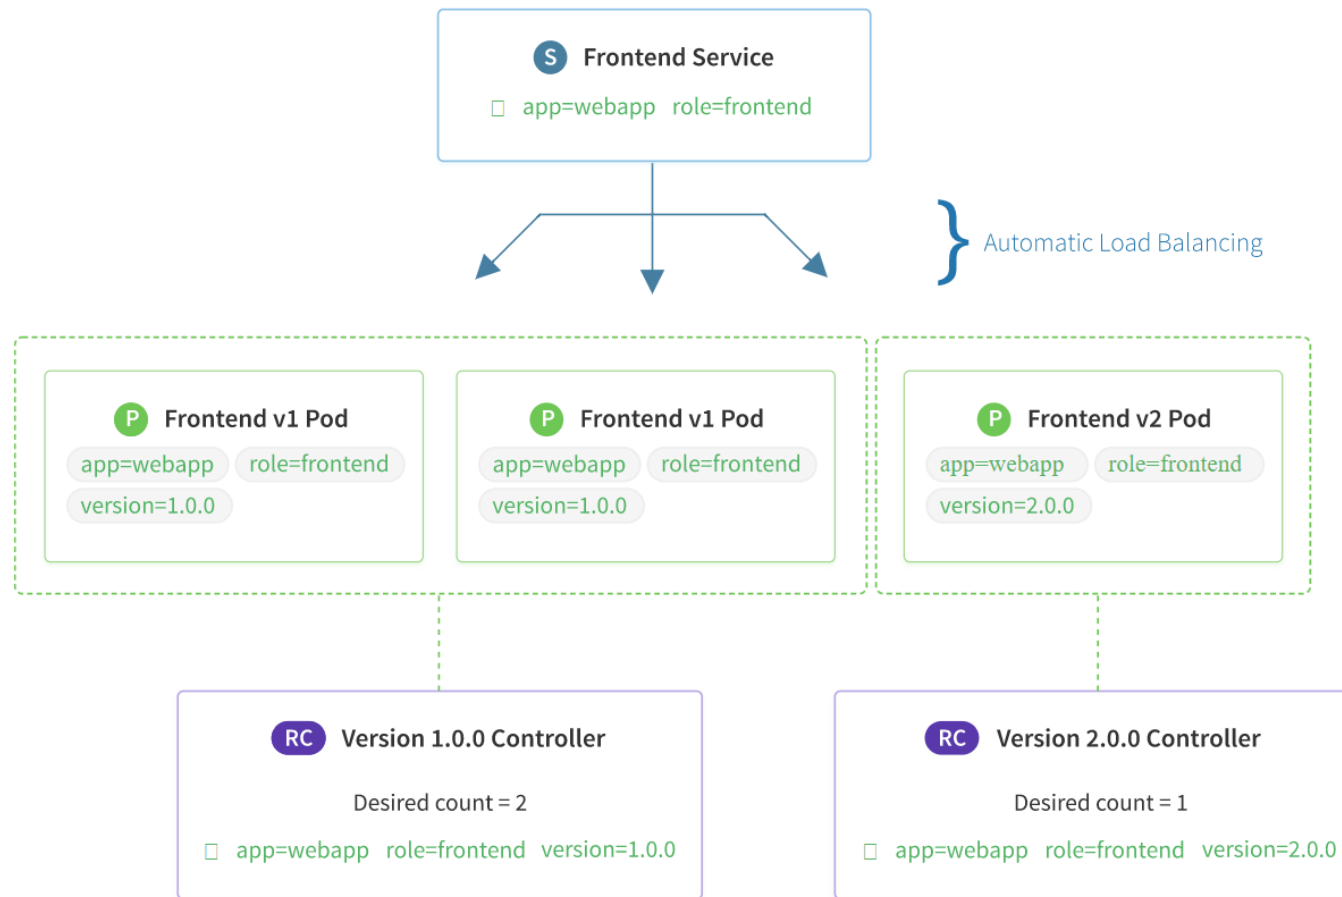


# Services – Load balancing

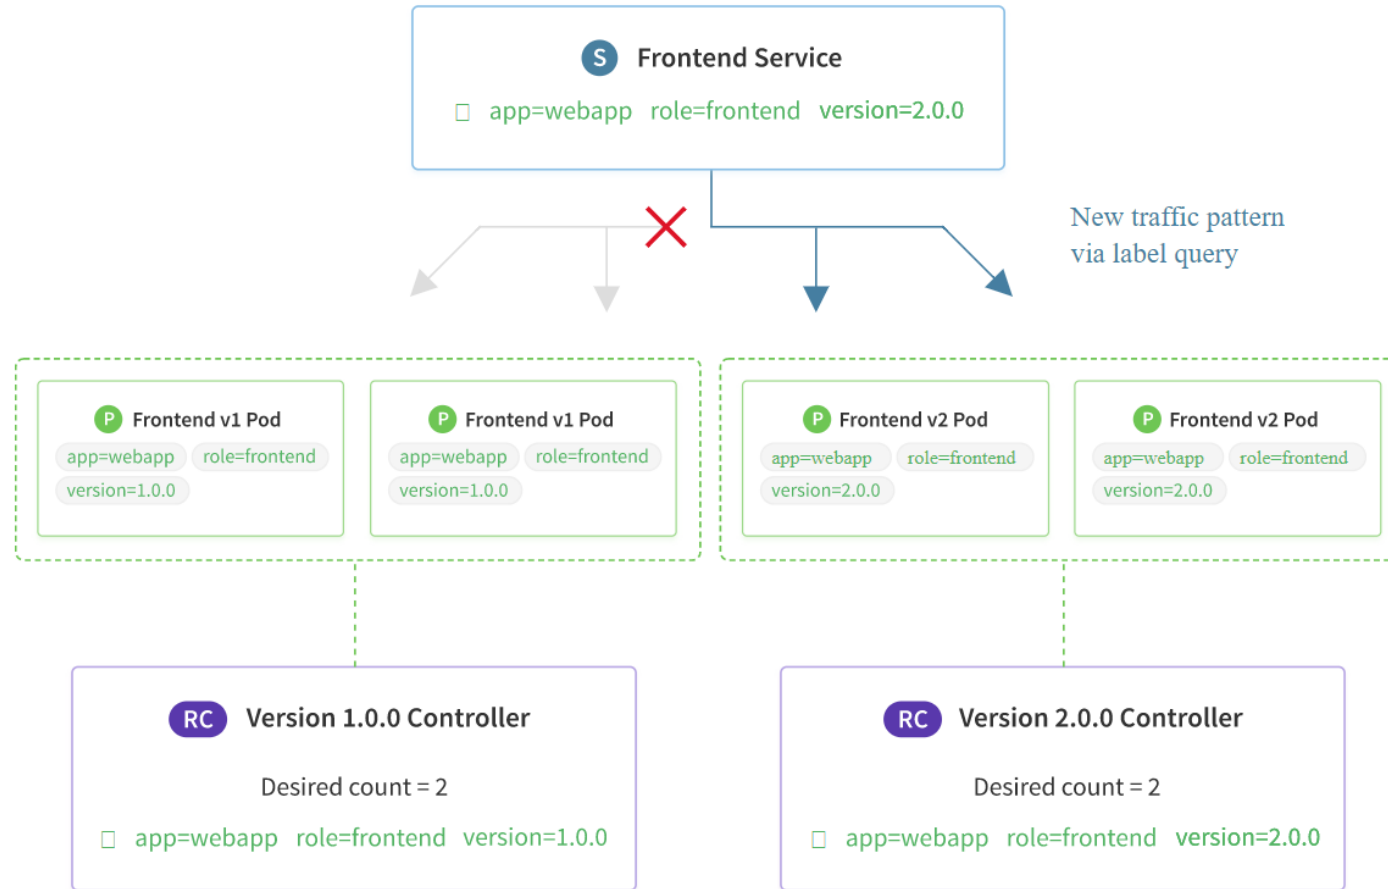
---



# Services - Rolling Deployment



# Services - Traffic Shift Deployment




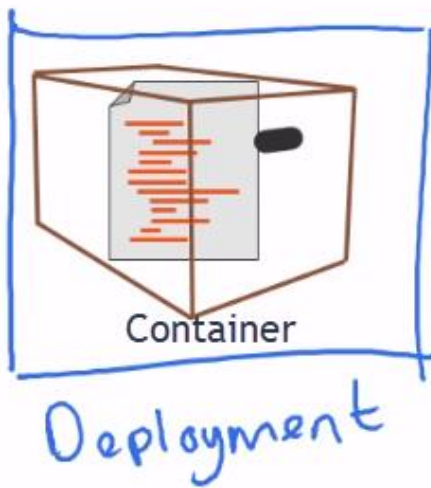
# Deployments

---

- Deployment controller provides declarative updates for Pods and ReplicaSets
- Describe the desired state in deployment configuration.
- Deployment controller changes the actual state to the desired state at a controlled rate.

# Deployments

---



# Secrets

---

- Secrets are intended to hold sensitive information, such as passwords, OAuth tokens, and ssh keys.
- Putting this information in a secret is safer and more flexible than putting it verbatim in a pod definition or in a docker image.
- Suitable while accessing private repositories – store repository login credentials.

# Create secrets

---

```
kubectl create secret docker-registry NAME --docker-username=username --docker-password=password --docker-email=email [--docker-server=string]
```

## Example

```
kubectl create secret docker-registry mysamplekey --docker-server=synacr.azurecr.io
```

```
--docker-username=acruuser --docker-password = 5//b=aG3Oo+/+ZBQhQSXCc=FEHA/Fv9w
```

```
--docker-email = username@mail.com
```

# Install, Configure Kubernetes

---



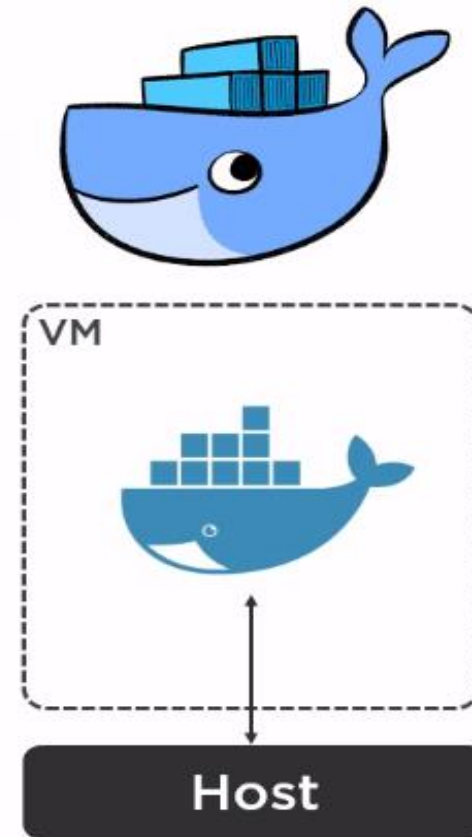
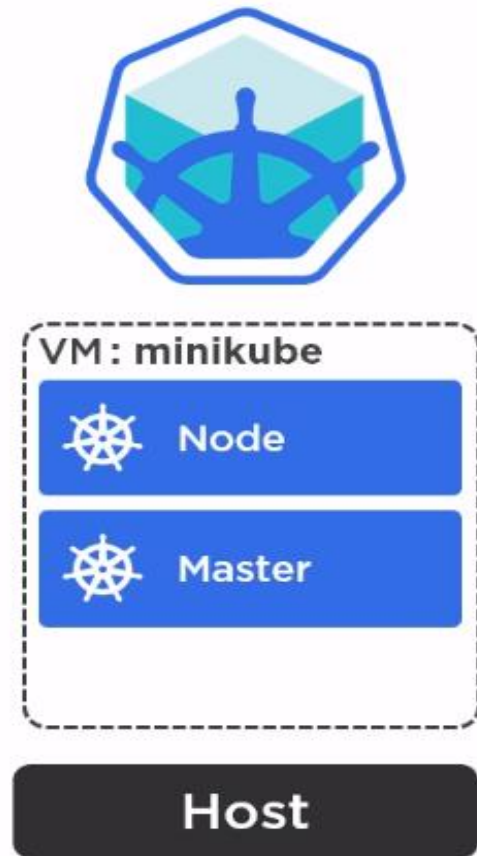
# Kubernetes

---

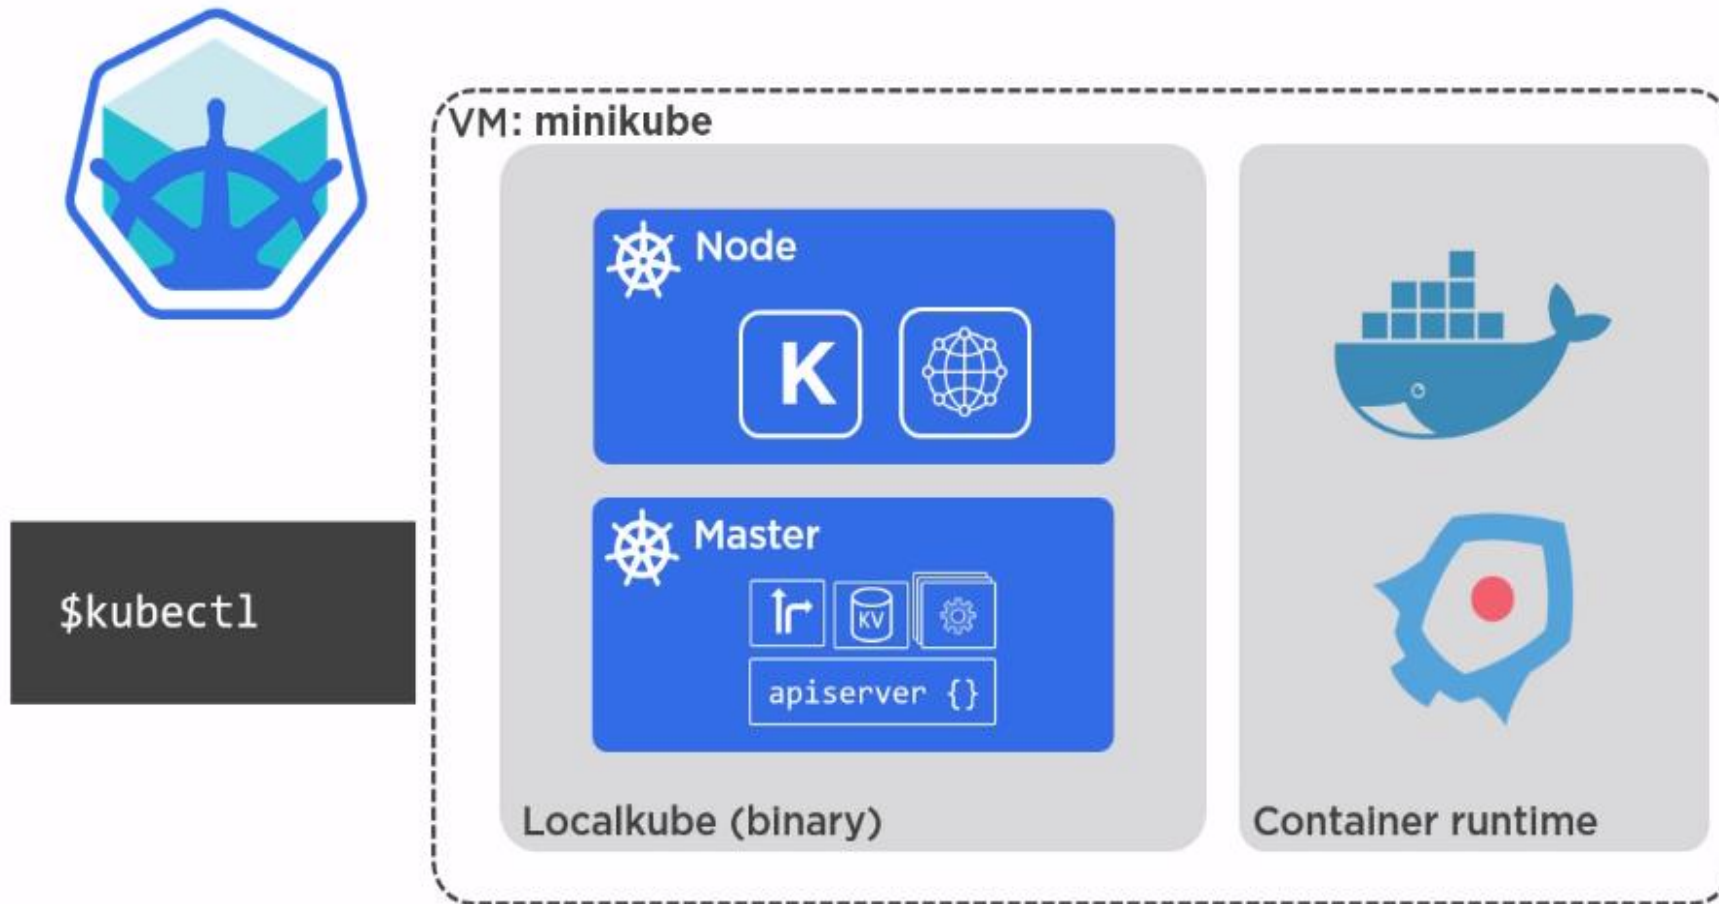
- Kubernetes can be installed and configured to manage clusters that runs Local machine, Cloud or on-premise.
- Local machine
  - Minikube - Tool that makes it to run Kubernetes locally.
- Cloud
  - Google Container Engine (GKE)
  - Azure Container Services (ACS), Azure Container Services – AKS

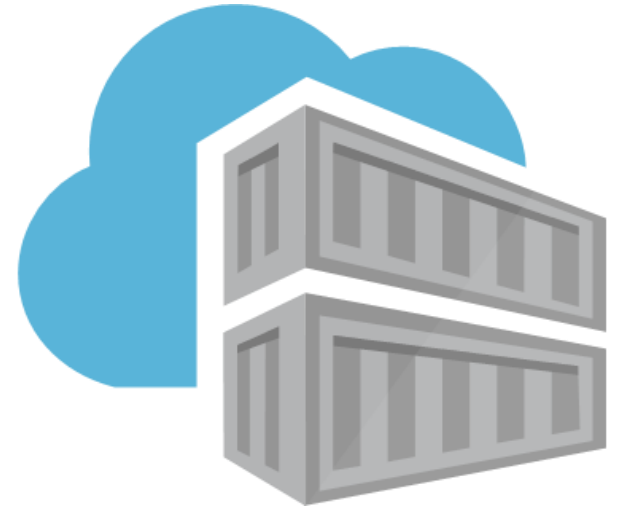
# Minikube cluster Vs Docker

---



# Minikube Architecture

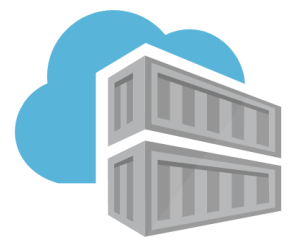




# Azure Container Registry

---

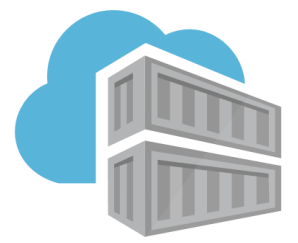
CONTAINER REPOSITORY SERVICE ON AZURE



# Azure Container Registry

---

- Private registry for hosting container images on Azure.
- Used for all container deployment Services
  - Azure WebApp for Containers
  - Azure Container Instance
  - Azure Container Services
- Supported by orchestrators – Kubernetes, Swarm, DCOS
- Push image using Docker Push
- Maintain windows and Linux container images
- Access management with AAD – secrets
- SKU – Basic, Standard, Premium



# Push Docker images to ACR

---

Login to Azure using Azure CLI

```
az login
```

Create ACR in resource group

```
az acr create --resource-group "ResGroup" --name "ACRName" -- sku basic
```

Enable admin login

```
az acr update -n "ACRName" --admin-enabled true
```

Login to ACR to run commands


```
az acr login --name "ACRName"
```

Tag Docker image with ACR login server name

```
docker tag <dockernamespace/imagename:tag> <ACR-loginname/imagename:tag>
```

Upload image to ACR

```
docker push ACR-loginname/imagename:tag
```



# Working with ACR

---

Login to ACR to execute Azure CLI commands

List repositories in ACR

```
az acr repository list -n <acrname> -o table
```

Display the login server name

```
az acr show --name <acrname> --query loginServer --output table
```

Display the ACR credentials

```
az acr credential show --name <acrname> --query "passwords[0].value"
```

# DEMO - ACR

---

DEPLOYING CONTAINERS USING WEB APP FOR CONTAINERS AND  
AZURE CONTAINER INSTANCE



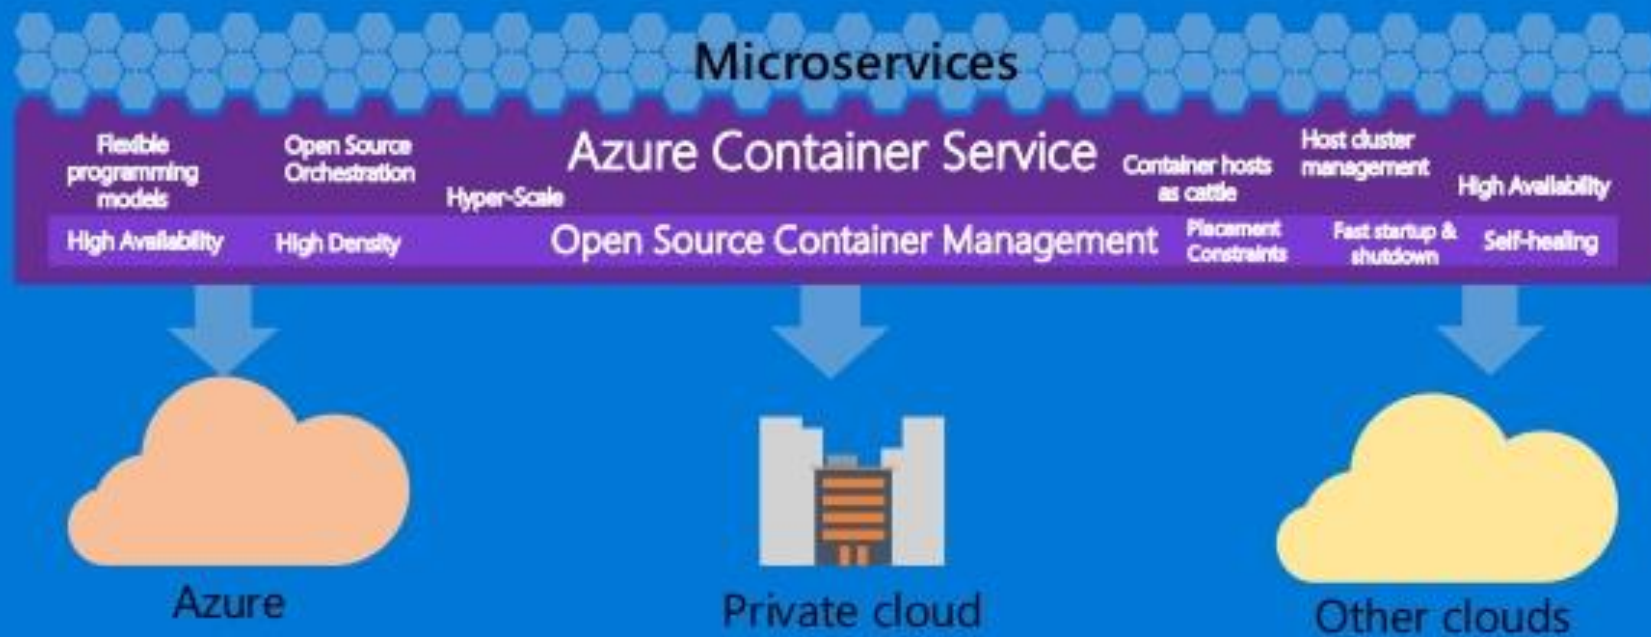


# Azure Container Services (AKS)

---

# Azure Container Service

A platform for reliable, hyperscale, container-based applications

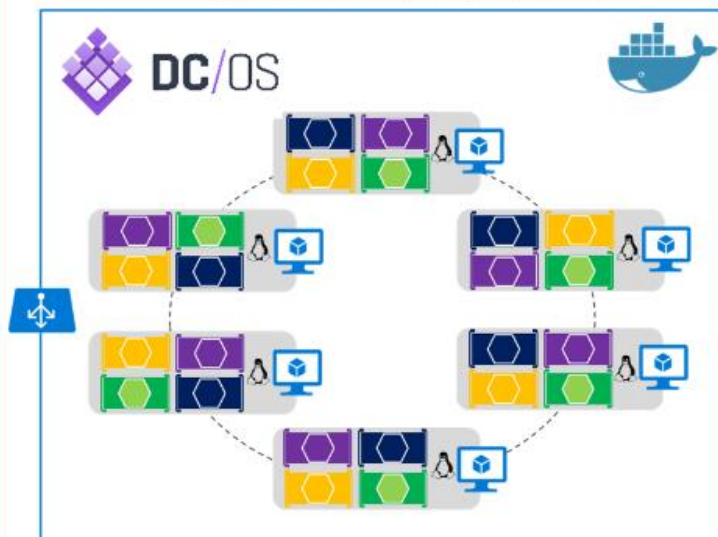





# Orchestrators for ACS

## Azure Container Service (ACS)


**Mesos DC/OS** cluster



**Kubernetes** cluster



**Docker Swarm** cluster





# Azure Container Services (AKS)

---

- Managed hosted environment for containers
- Kubernetes clusters
- On demand provisioning, upgrading and scaling
- Health monitoring and maintenance by Azure
- Pay only for Agent nodes not for Masters

# AKS

---

## AKS Offers

- Automated Kubernetes version upgrades and patching
- Easy cluster scaling
- Self-healing hosted control plane (masters)
- Cost savings - pay only for running agent pool nodes
- Automatic cluster upgrades

# AKS using Azure CLI

---

Enabling AKS preview for your Azure subscription

```
az provider register -n Microsoft.ContainerService
```

Creates Kubernetes cluster

```
az aks create --resource-group myResourceGroup --name myK8sCluster \  
--node-count 1 --generate-ssh-keys
```

Connect to the cluster

```
az aks install-cli
```

```
az aks get-credentials --resource-group myResourceGroup \  
--name myK8sCluster
```

```
kubectl get nodes
```

# Scaling Nodes

---

```
az aks scale --resource-group=<grp-name> \  
  --name=<AKS name> \  
  --node-count 5
```

# Scaling pods

---

Manually scale pods

```
kubectl scale --replicas=<pod-count> deployment/<deployment-name>
```

Auto-scaling pods

```
kubectl autoscale deployment <deployment-name> \  
--cpu-percent=50 --min=3 --max=10
```

List autoscaler status

```
kubectl get hpa
```



# Update deployments

---

Update deployments with new version of application image

```
kubectl set image deployment <deployment-name> \  
    <container-name>=<updated-image-name>
```

# Upgrade Kubernetes in Azure Container Service (AKS)

---

Get cluster versions

```
az aks get-versions --name <aks-name> \
    --resource-group <group-name> --output table
```

Name	ResourceGroup	MasterVersion	MasterUpgrades	NodePoolVersion	NodePoolUpgrades
default	DemoGroup	1.7.7	1.7.9, 1.8.1, 1.8.2	1.7.7	1.7.9, 1.8.1, 1.8.2

Upgrade cluster

```
az aks upgrade --name <aks-name> --resource-group <group-name> --kubernetes-version
<desired version>
```

Eg: az aks upgrade --name bst-cluster --resource-group DemoGroup --kubernetes-version 1.8.2

# Thank You

---



sonusathyadas@hotmail.com



<https://in.linkedin.com/in/sonusathyadas>



@sonusathyadas